ABSTRACT

Title of dissertation: ADDING GRADUAL TYPES
TO EXISTING LANGUAGES:
A SURVEY

Type system defines sets of rules, which is checked against programs to prevent certain kind of errors and bugs. Type checking happens either statically ahead of execution, or dynamically at runtime. Programming languages make different choices regarding type system to accommodate their specific needs.

Statically typed languages are reliable and efficient. Because type errors are caught ahead of execution, programs are free of these errors at runtime. In addition, having type information available statically also helps development in many other ways beyond the type system, such as type-specialized optimization and improved development tools. While some believes static type systems assist code evolution in long run, a new feature or design change in codebase could introduce great effort to not just implementation but also adjusting type annotations accordingly, which might not be desirable for testing out new ideas on top of an existing large codebase.

Dynamically typed languages stand out when it comes to scripting and fast prototyping: executables from external sources can be used directly without worrying about type information and programs can be easily modified and executed without need of enforcing consistency and correctness throughout the program. But a program grows over time to handle more complicated tasks, the interaction between

different portions of the program becomes harder to track and maintain.

It is only natural that we start exploring the spectrum between them: type systems that enjoy benefits of both. Among one of these is gradual typing, which features a type system that provides optional type checking. In such a system, programmer can give types to only portion of the program. For the portion that has type information, static and dynamic typechecking will work together to ensure type consistency, whereas the untyped portion of the program is left unchecked, which might be desirable when performance or flexibility is concerned. By giving programmers control over when and where should typechecking occur, one does not commit to either static or dynamic typing while losing the benefit of the other.

One crucial advantage of gradual typing is that it can be implemented by extending an existing statically or dynamically typed language. Syntactically, the extension often results in a superset of the original language that allows more expressiveness in types, which means the effort of migrating existing code and language users to take advantage of gradual typing is minimized. Semantically, adding or removing type annotation does not change the result of programs aside from type errors, which grants gradual and smooth transition between static and dynamic type disciplines.

In this survey, we will walk through the relevant research literature on extending existing languages to support gradual types, discuss challenges and solutions about making these extensions, and look into related works and the future of gradual typing.

Table of Contents

Chapter 1:   Introduction

Type system defines sets of rules about programs, which can be checked by a machine in a systematic way. One can assign types to appropriate language concepts like variables, functions and classes to describe expected behaviors and how they interact with each other. Then type system takes the responsibility of ensuring that types are respected and type errors prevent erroneous part of a program from execution.

There are other benefits of type systems besides automated checking: types can be used to reveal optimization opportunities, can serve as simple documentation to programmers or provide hints to development tools to automate navigation, documentation lookup or auto-completion.

All these benefits make type system an important part of programming languages, and there are different ways of accommodating languages with type systems. The most noticeable difference is the process of checking programs again type rules, which is also known as type checking. It occurs either statically ahead of program execution, or dynamically at runtime. For a static type system, ill-typed programs are rejected by compiler and no executable can be produced until all type errors are fixed. For a dynamic type system, type information is examined at runtime and

type errors result in abortion of the program or exceptions being raised instead of causing segmentation faults or other worse consequences. Static and dynamic type system both have their advantages and disadvantages, and languages make different choices depending on their needs.

In this chapter, we start off by discussing static and dynamic type systems, which motivates research works that attempted to combine both within one system in order to take the benefits from the two, which in return gives birth to many possible solutions. One particular example among them is gradual typing.

1.1   Strengths and Weaknesses of Static and Dynamic Typing

The difference between static and dynamic typing is the difference between whether type checking occurs ahead of execution or at runtime. This section discusses advantages and disadvantages of both.

In this section we use $\tau_0, \tau_1, \ldots$ for type variables. **bool**, **num** and **str** are types of boolean values, numbers and strings respectively. Tuple types are notated as $(\tau_0, \tau_1, \ldots)$, and the type notation for functions that takes as input a value of type τ_0 and returns a value of τ_1 is $\tau_0 \rightarrow \tau_1$. We use $o : \tau$ to mean that a term o is assigned type τ. For example $f : (\textbf{num}, \textbf{num}) \rightarrow \textbf{num}$ is a function that takes a tuple of two numbers as input and returns a number.

1.1.1 Static Type Systems

A static type system is preferred if robustness and performance is the main goal of a language. Because type checking occurs ahead of execution, ill-typed programs are rejected before it can start, and when a program typechecks, one can therefore expect no type error to occur and no extra cost is paid for typechecking at runtime.

For a static type system to work, all variables, functions, etc. in a program will need to be assigned types. This is usually done by programmers or through the means of type inference, which is a technique that infers types using available type information. This is both an advantage and a disadvantage of a static type system: having type annotations improves readability and since programmers are required to keep the consistency between type and code, type also serves as simple, faithful documentation. But on the other hand, adding and maintaining type annotations can also considered a burden.

Having static known type information also helps in terms of performance in other ways. For example, In many programming languages, arithmetic functions are polymorphic. It is allowed for $\mathbf{a}$ and $\mathbf{b}$ to have different numeric types in $\mathbf{a} + \mathbf{b}$, a cast will be inserted to ensure arithmetic primitives only deals with addition of compatible types: if we are adding integer $\mathbf{b}$ to a floating number $\mathbf{a}$, then $\mathbf{b}$ will be casted so we can call addition primitive that expected floating numbers on both sides. A dynamic language would have to figure out types of $\mathbf{a}$ and $\mathbf{b}$ at runtime, make decision about whether casting is required or which primitive addition to call all at runtime. But with type information statically available, these decisions can

be made ahead of execution, resulting in improved performance at runtime.

Besides extra effort of maintaining type annotations, the disadvantage of a statically type language often lies in the lack of runtime flexibility. If we want to write a program that deals with data whose structure is unknown at compile time, runtime inspection must be possible. While typical dynamically typed languages allow inspect of types or object properties, some extra work are required for a statically typed language to keep type information available at runtime.

In addition, it is often less convenient for programmers to prototype in statically typed languages: it involves trial and error to find the best implementation to solve a problem, which means the ability to write partial programs, make frequent changes to structure of variables and functions, but these features are not easily available for a statically typed language.

1.1.2   Dynamic Type Systems

Dynamic type systems do typechecking at runtime, which is best suited for languages designed for scripting and fast prototyping. A shell script, for example, runs other executables and feeds one's output to another, effectively gluing programs together. In such a case, we have little to no knowledge about these executables ahead of time, making statically type checking impractical. And when it comes to prototyping, it is more important for programmers to execute the code and make modifications accordingly than to enforce correctness and consistency through whole program. In such scenario, it is more convenient to delay type checking until it is

required at runtime.

However, it is also easy to spot weaknesses of dynamic type systems. Despite that type checking does not happen ahead of execution, programmers usually have facts about behaviors of programs, which is often described in comments or through naming of variables and functions without a systematic way of verifying them. As program develops, it is easy to make changes and forget updating these descriptions accordingly, which forces future maintainers to go back and rediscover these facts, causing maintenance difficulties.

Furthermore, without type information statically available, dynamic type system needs to maintain runtime information and rely on it to implement expected semantics. For example, if it is allowed for a language to condition on non-boolean values, its type has to be available at runtime in order to determine how to cast it into a boolean value at runtime.

Optimization could also be hindered for a dynamically type language: array cloning can be implemented efficiently as memory copy instructions knowing the array in question contains only primitive values but no reference or pointers, whereas without prior knowledge like this, every element of the array has to be traversed and even recursively cloned.

1.2   Combining the benefits of the two

One might have noticed that static type checking and dynamic one are not mutually exclusive: it is possible to have a program typechecked statically, while

allowing runtime type checking. While static type system does typechecking ahead of time and avoids runtime overhead, it lacks the ability of using runtime type information; dynamic type system does typechecking at runtime, but it requires sophisticated work to reduce the amount of unnecessarily repeated runtime checks. In hope of bringing in benefits of both, several attempts are motivated to put these two type systems in one. In this section, we will discuss about these research works.

1.2.1 Design Goals

Summarizing pros and cons of static and dynamic typed languages, we now have a picture of an ideal type system, it should have following features:

- **Detect and rule out errors ahead of program execution.** This is one main purpose of having a static type system. When a program starts running, type errors should already be eliminated and runtime overhead should be minimized to keep its performance close to that of a static type system.

- **Delay typechecking until needed at runtime.** This is the crucial feature of a dynamic type system, it allows ill-typed or partially written programs to execute, which makes it suitable in situations where static type information is limited or fast prototyping is preferred over robustness or runtime performance.

Aside from these two goals, type annotations can also be used for other purposes like optimization and providing hints to development tools. For a language with this new type system, it should also be possible.

1.2.2   A Introduction to Gradual Typing

Gradual typing is one possible approach that meets our design goal. It captures the fact that type information might only be partially known ahead of execution by introducing a special type that indicates partial types. It attempts to typecheck programs like a static type system does, but delays typechecking on partial types until sufficient type information is available at runtime.

In the following sections, we will introduce it as an extension to static type system. But it is also possible and practical to support gradual typing by extending a dynamic type system.

1.2.2.1   Dynamic Type and Type Consistency

Gradual type system introduces a special type on top of a static type system to capture the idea that type information might only be partially known ahead of execution. The special type is usually called a "dynamic type", we introduce a special type **dyn** for it.

To deal with dynamic types, type consistency is introduced to accommodate type judgments for gradual typing. Figure ? shows rules for type consistency relation.

(TODO: type consistency rules here, lambda calculus and with some ground types)

In general, two types are consistent when their type structures do not conflict with each other. For example, **bool** is consistent with type **dyn** despite that the

latter is less precise; **num** → **dyn** is consistent with both **dyn** → **num** and **dyn** → **dyn**, because all these types are functions that accepts one value and returns another and **num** and **dyn** are consistent with each other. However **dyn** → **dyn** is never consistent with **bool**, because the former is a function type while the latter a boolean value, whose type structures do not match.

In order to perform static typechecking, it is required that all terms in the program to must have types ahead of execution. In a gradual type system, this is no longer required: for a specific term, if programmers choose not to give a type annotation, it will be assigned type **dyn**. This still allows static typechecking to detect errors because for types that contains **dyn**, it is not totally ignored, but checked up to the point where the structure of type is known.

On the other hand, a dynamically type program can now be statically type-checked by considering it to be a program without any type annotation. This is still beneficial, as some literal values and primitives will still have their built-in types, an obvious misuse like taking the square root of a string value can now be detected ahead of execution instead of raising runtime exceptions.

1.2.2.2   Cast Insertion and Runtime Typechecking

As we can see gradual type system attempts to contain both statically and dynamically typed programs. However, because of presence of dynamic types, the type system cannot be sound by only have static typechecking. Imagine function $f : $ **num** → **dyn** and variable $v : $ **dyn**, while the function application $f v$ can pass

static typechecking, the runtime value of t might turn out to be a string and now the application is ill-typed.

The problem lies in the fact that some terms might be assigned with types that contains **dyn**, its runtime counterpart are precise, which is not checked during static typechecking. To solve this problem, casts are inserted at appropriate locations, which then allows type errors to be detected at runtime. Suppose we have a cast function $c : \mathbf{dyn} \to \mathbf{num}$, function application fv then becomes $f(cv)$, which still typechecks, but whenever we need to pass a runtime value to f, its runtime is always checked by c to decide whether to pass the value to f or raise a type error.

1.2.2.3 Optional Type Annotation and "pay-as-you-go"

So far we have seen that gradual type system introduces a dynamic type, which allows accepting both statically and dynamically typed programs during static typechecking. But what is more important is that it creates a middle ground in which programs can be partially annotated with types: in a gradual type system, terms are either explicitly given a type, or, when type annotation is missing, given **dyn** type. This renders all type annotations to be optional and allows programmers to have more control over what they need from typechecking: for an originally statically typed program, we can now write terms without giving a type, this allows easy prototyping; for an originally dynamically typed program, some parts of it might already been stable and robust enough, annotating these parts with type allows some effort of ensuring correctness being shifted to typechecking.

9

For example, one might wish to implement a function f that expects a string value as input. There are two approaches: we can either give f a type annotation: $f : \mathbf{str} \to \mathbf{dyn}$. By doing so, type system takes the responsibility of making sure that the input type is indeed a string. Alternatively, we can inspect type of the input value at runtime ourselves. A static type system allows us to do the former, but a return type must either be given explicitly or inferred. And a dynamic type system does the later, but we cannot statically detect any mistake like applying a number literal to f. However, a gradual type system finds the balance between two, it allows us to give input value a type without worrying about giving a return type if not needed, therefore programmer only pays for what they want from a type system.

1.2.2.4 Gradual Guarantee

Regarding the term gradual typing itself, Siek, Vitousek, Cimini and Boyland's work gives a detailed explanation of the original intention. Several works we are going to discuss in the following few chapters respects such idea:

- Gradual typing includes both fully static and fully dynamic

- Gradual typing provides sound interoperability

- Gradual typing enables gradual evolution

1.2.3 Alternatives

There are other type systems that bring benefits of the two together. Cartwright and Fagan's introduced soft typing, which is a type system that programers do not write type annotations but use type inference to assign appropriate types to terms, and runtime checks are inserted as needed. The drawback of this design is that programmers do not have control over types. Quasi-static typing is another attempt, with the idea of dynamic type in use, it chooses to use a subtyping relation and allows both up-casts and down-casts. A second pass of plausibility checking detects incompatible casts and signals the program in question as ill-typed. However its typechecking algorithm did not receive a correct proof, and it does not statically catch all type errors (TODO: ref) in the process of attempting the proof, Siek and Taha finds a better solution, which results in gradual typing that we have seen today.

1.2.4 Discussion

Among various research works that combines benefit of static and dynamic type checking gradual typing remains one of the most popular approach. In the sprite of gradual guarantee, gradual typing extensions aim at making all existing programs written in the original one still valid with little difference in semantics. In addition, unlike other approaches, programmers have better control over should and should not be typechecked.

Despite that gradual type systems are well-studied research topics, extending existing languages to support them is far from trivial task. To name few differences,

while theoretical type systems assume primitive types as simple as just numbers and booleans, a practical language have various kinds of them including integers, floating numbers, strings and arrays, which have not been addressed much. In addition, language-specific features, idioms and common practice needs to be carefully studied to allow existing code and language users to transit smoothly into the extended language. In this survey we will explore literatures that extends existing languages to support gradual typing. In chapter 2, we will focus on benefits of making these extensions. Chapter 3 discusses and categorizes common and language-specific challenges researchers have to face and solutions to them. Some of these challenges, however, still remains as open questions, which will be covered in chapter 4.

Chapter 2: Benefits of Gradual Typing

Gradual typing brings existing languages not just the good parts of both static and dynamic typing. The synergy also creates surprising new benefits to languages users.

In this chapter, we will explore some literatures and see in detail about these benefits that gradual typing have brought in.

2.1 From JavaScript to Safe TypeScript

Starting as a scripting language, JavaScript has grown into a language that powers many large web applications. Unfortunately, as a language that exists over a long period of time, the language evolves but many of its flaws are still left as it is for compatibility reasons, hindering productivity of many programmers.

Among tools and extensions to JavaScript that attempt to ease this pain, TypeScript is one closely related to gradual typing: it is a superset of JavaScript that supports optional type annotations, a compiler typechecks TypeScript code ahead of execution and compile code into plain JavaScript source code, making it ready to be executed out of box for JavaScript interpreters. However, TypeScript is intentionally unsound: only static typechecking is performed, which does not

prevent runtime type errors. One noticeable improvement to TypeScript is Safe TypeScript by Rastogi, Swamy, Fournet, Bierman & Vekris, which shares the same syntax with TypeScript but features a sound type system and efficient runtime type information (RTTI)-based gradual typing.

In this section we will visit some language features of Safe TypeScript and see how it improves JavaScript. Note that TypeScript uses **any** for dynamic type.

2.1.1 Type Checking for Free

Despite JavaScript is capable of examining type of values at runtime, it does not have support for type annotations. This leads to a common practice in which function bodies will have code just for checking its argument types and throw errors before proceeding with actual implementation:

```
function f(x) {

    if (typeof x !== 'string') {

        // throw error

    }

    return x + '!';

}
```

As a toy example, the function accepts a string value, then returns another with exclamation sign concatenated to it. But imagine in real projects there will be multiple arguments to a function and some of them have to go through this process of checking, it will soon become less maintainable.

By extending the language with type annotations, we can do something better in Safe TypeScript:

```
function f(x : string) : string {

    // now this check becomes unnecessary

    if (typeof x !== 'string') {

        // throw error

    }

    return x + '!';

}
```

In the code above, we just declared a function **f** that accepts one variable x of **string** type, and returns a value of **string** type. The code for runtime checking is removed, instead, Safe TypeScript performs typechecking and insert casts when needed to ensure that, when **f** is called, its argument is indeed a **string**. This renders the **typeof** check at the beginning useless: when we enter the function body **x** is guaranteed to be a **string**, therefore the body of the **if** statement cannot possibly be entered. This allows us to get rid of the check totally:

```
function f(x : string) : string {

    return x + '!';

}
```

The use site of **f** might look like the following:

```
f('one'); // good
```

```
f(1); // bad

function g(x) {

    return f(x); // might be unsafe

}
```

Here Safe TypeScript is able to tell that: calling **f** with a string literal is safe, and no extra cast is needed; the second call is immediately rejected because number literal is clearly not a **string**; and for the third case where **x** is not explicitly given a type, Safe TypeScript compiler will insert runtime type cast and either allow it to proceed to the body **f** when **x** is indeed a string, or throw an error before even calling **f**.

Notice that by making use of type annotations, Safe TypeScript not only allows clearer code, but also able to spot some type errors ahead of execution: imagine the function application **f(1)**, when its written in plain JavaScript, we have to wait for the check in function body to throw an error, while annotating **x** with a type allows Safe TypeScript to spot the error even before execution.

2.1.2   Object-Oriented Programming in Safe TypeScript

JavaScript is a prototype-based language. This means runtime objects form a chain in which one object can be a prototype of other objects and to perform method invocation, method names are searched first in object itself and then along the chain, and the first one with matching method name is used for the invocation. Using well-known techniques allows JavaScript to simulate objected-oriented programming.

Designed to be a superset of JavaScript, Safe TypeScript must accommodate such programming style as well.

TypeScript introduces the concept of interface to allow representing the structure of runtime objects.

```
interface Point { x: number; y: number }
```

This defines **Point** type to be an object that has two fields **x** and **y** whose types are all **number**s. For example, the invocation with object literal $f(\{\mathbf{x: 1, y: 0}\})$ typechecks with $f : Point \Rightarrow any$. An interface can be implemented by classes:

```
class MovablePoint implements Point {
  constructor(public x: number, public y: number) {this.x = x; this.y = y;}
  public move(dx: number, dy: number} { this.x += dx; this.y += dy; }
}
```

While in TypeScript, all classes are viewed structurally, this does not match with JavaScript semantics. Consider the following function:

```
function mustBeTrue(x : MovablePoint) {
  return !x || x instanceof MovablePoint;
}
```

One might expect this function to always return **true**. However, a structural view will consider object literal $v = \{\mathbf{x: 0, y: 0, move(dx: number, dy:number)}\}$ to match with **MovablePoint**, but **mustBeTrue(v)** will return **false** because

17

JavaScript will expect v to have **MovablePoint** somewhere along its prototype chain.

To solve this problem, Safe TypeScript takes a different approach: it treats class types nominally but can be viewed structurally. Namely, **MovablePoint** is still a subtype of $t_m = $ **{x: number, y: number, move(dx: number, dy:number): any}** and **{x: number, y: number}**. But t_m is no longer a subtype of **MovablePoint**.

2.2   From Scheme to Typed Scheme

Scheme is a dynamically typed language. It is used for casual scripting as well as industrial applications. Like other dynamic typed languages, programmers are not attached to using any type discipline but rely on various kind of reasoning on their needs. This makes the language flexible but challenging to assign proper and precise types to terms.

Tobin-Hochstadt and Felleisen's work on Typed Scheme introduces the notion of occurrence typing, which extends Scheme with gradual typing.

2.2.1   Support Informal Reasoning

The following code shows a function definition as typical style of programming in this language:

```
;; a Complex is either

;; - a number
```

```
;; - (cons number number)
(define (creal x)
  (cond [(number? x) x]
        [else (car x)]))
```

The code above defines a function **creal**, which takes as argument a complex number: a real number is simply a value satisfying **number?**, while an imaginary number is a pair of **number**. Function **number?** distinguishes two different representation of numbers: in the the first branch of **cond**, **x** is treated like a number while in the second branch **x** is a pair and **car** is used to extract its real part.

One might have noticed that while input value **x** could be either a number or a pair of numbers, predicate **number?** distinguishes between them and in difference branches of the **cond** expression, **x** gets different types.

It is important that type system should be able to assign same variable with different types depending on the context where variable occurs, this is exactly the notion of occurrence typing.

The Typed Scheme version begins with definition of complex number:

```
;; a Complex is either
;; - a number
;; - (cons number number)
(define-type-alias Cplx (Union Number (cons Number Number)))
```

While **Cplx** is just an alias for the union type, it improves readability and allows other parts of the program to refer to it by name. The body of the function

looks like:

```
(define: (creal [x: Cplx]) : Number
  (cond [(number? x) x]
        [else (car x)]))
```

Note that we are explicitly giving the return type **Number**, and the type system is capable of inferencing that both branch of **cond** expression returns a number so it will still typecheck.

Another different kind of reasoning is also being used among Scheme programmers: Suppose we have stored a list of various types of values in **xs** and the following code will compute the sum of all numbers from **xs**:

```
(foldl + 0 (filter number? xs))
```

Note that despite **xs** is a list that contains not just numbers, it is guaranteed that **(filter number? xs)** will return a list of numbers so **foldl** will work properly to produce the desired result.

The expression requires no modification at all to typecheck in Typed Scheme and it receives type system benefits: suppose **map** is mistakenly used instead of **filter** or **number?** is replaced by other predicates insufficient to test whether a value in question is indeed a value, the type system is able to recognize such errors and give warnings.

2.2.2   Refinement Type

Note that in Scheme, we use **number?** to distinguish numbers from other type of values. In general, given a boolean-valued unary function we can define the set of values that produces **true** when fed to this function. This allows Typed Scheme to support refinement type.

```
(: just-even (Number -> (Refinement even?)))
(define (just-even n)
  (if (even? n) n (error 'not-even)))
```

The code above uses boolean-valued function **even?** and its corresponding refinement type **Refinement even?**, so we can expect a value of such type to be not just numbers, but also only even ones.

More practical examples include use refinement types to distinguish raw input from validated ones:

```
(: sql-safe? (String -> Boolean))
(define (sql-safe? s) ...)
(declare-refinement sql-safe?)
```

In this example, user defines function **sql-safe?** to verify that a raw string contains no SQL injection or other contents of malicious attempt. Then a refinement type is declared using this very function. This makes available type **(Refinement sql-safe?)** to rest part of the program to avoid mistakenly using raw input rather than verified safe ones.

As we can see occurrence typing not just enables gradual typing, but also allows refinement type, which provides us a powerful tool to use user-defined functions to define custom types of better precision.

2.3   From Python 3 to Reticulated Python

Python 3 (Python for short) is another popular dynamic typed language. M. Vitousek, M. Kent and Baker's work on Reticulated Python (Reticulated for short) explored extending Python with gradual typing.

Python comes with annotation syntax. This syntax allows expressions to be optionally attached to function definitions and their arguments, which makes it suitable for type annotation.

Given proper definition of **Int**, we can write annotated function **distance** like below:

```
def distance(x: Int, y: Int)-> Int:
  return abs(x - y)
```

While annotation allows arbitrary expressions, Reticulated Python only uses types built from several type constructors.

2.3.1   Function Types

Python 3 has different ways of calling the same function. Suppose we want to use **distance** defined above, we can use positional arguments like **distance(3,4)**, or use keywords like **distance(y=4, x=3)** and the result should be the same.

To support keyword calls, **Named** constructor is used to give **distance** type **Function(Named(x: Int, y: Int), Int)**. To support traditional positional arguments, **Pos** is used instead. The type system also allows **Named** to be used as if it is constructed by **Pos** when their length and element types correspond. So functions can be called in both ways without problem. Additionally **Arb** is another constructor that allows functions of arbitrary parameters.

2.3.2 Class and Object Types

Python is an object-oriented programming language. Programmers define classes and create objects from them. Both classes and objects are runtime values in Python and Reticulated provides corresponding types for both.

Consider the following example:

```
@fields({'x': Int})

class 1DPoint:

  def __init__(self: 1DPoint):

    self.x = 0

  def move(self: 1DPoint, x: Int)->1DPoint:

    self.x += x

    return self


p = 1DPoint()
```

It defines a class **1DPoint** with one field **x** of **Int** type. And an object is

23

created from it through the constructor and is bound to variable **p**.

Reticulated derives class type from this definition:

```
Class(1DPoint){
  move: Function(Named(self: 1DPoint, x: Int), 1DPoint)
}
```

And object types are similar:

```
Object(1DPoint){
  move: Function(Named(x: Int), 1DPoint)
}
```

While object can access methods of its class, it is also possible to assign methods and fields to objects to change its behavior. Therefore Reticulated makes object type open with respect to consistency.

2.3.3 Dynamic Semantics

Reticulated Python is not only a practical gradual type extension, but serves as an experiment of exploring different dynamic semantics

2.3.4 Other Features

There are other features of Reticulate. In particular, since it is not always possible to know what module would be loaded at runtime, static typechecking sometimes occurs right after module is loaded.

Despite Python syntax supports annotation at function definition, it does not provide a way of annotating local variables with tyoe. Instead, Reticulated takes a step further and perform dataflow-based type inference.

2.4 From Smalltalk to Gradualtalk

Smalltalk is a language of highly dynamic nature: it supports live programming, which means programs with incomplete methods should be accepted, and programmers in Smalltalk relies on idioms that are tricky to type properly. Allende, Callau, Fabry, Tanter and Denker's work on Gradualtalk shows us a well combination of various features in order to extend Smalltalk into a gradual typed language.

2.4.1 Annotating Programs with Types

In Smalltalk there are objects that receives and processes messages. The following example computes euclidean distance for points:

```
Point >> distanceTo: p
| dx dy |
dx := self x - p x.
dy := self y - p y.
^ (dx squared + dy squared) sqrt
```

Gradualtalk extends this syntax that allows type annotation on parameter, return value and local variables.

```
Point >> (Number) distanceTo: (Point) p

| dx dy |

dx := self x - p x

dy := self y - p y

^ (dx squared + dy squared) sqrt
```

Since local variable **dx** and **dy** does not explicit type-annotated, they receive **Dyn** as default type.

Blocks are basic features in Smalltalk, whose corresponding types are available in Gradualtalk in the form of normal function types.

```
Polygon >> (Number) perimeter: (Point Point -> Number) metricBlock

...
```

This definition expects **metricBlock** to be block that takes as argument two points, and returns a number.

2.4.2   Self Types

One practice in Smalltalk is to return object itself for setters to allow chained calls:

```
Point >> (Point) y: (Number) aNumber

y := aNumber.
```

Notice that if we want to extend this class in future, its setter with be tagged with type **Point** and some type information will be lost in the way. To deal with this problem, Gradualtalk introduces self type:

```
Point >> (Self) y: (Number) aNumber

y := aNumber.
```

If any other class inherits from it, the self type will make sure to point to that class instead of sticking with **Point**, this allows type information to be preserved.

2.4.3 Union Types

Gradualtalk supports union types, which is useful when source code contains different branches and each of them might return different types.

```
Boolean >> (a | b) ifTrue: (-> a) trueBlock ifFalse: (-> b) falseBlock
```

Note that **trueBlock** and **falseBlock** might return different types, allowing this method to have `a` or `b` will not be satisfactory, and simply using `Dyn` type returns in a lose of information. Gradualtalk solves this problem by introducing union types: `a | b` is a type by itself, which is compatible with both `a` and `b`.

2.4.4 Structural and Nominal Types

A structural type describes the shape of an object.

```
RBParser >> bracketsOfNode: ({left (-> Integer) . right (->Integer)}) node
```

The method above defines the argument type of `node` to have 2 methods: `left` and `right` and both of them will return an integer when invoked.

A nominal type, on the other hand, is induced by classes (for example an instance of `String` will have `String` type). In addition, subtyping relations of them are formed from existing inheritance relation.

27

Gradual type managed to unify them in an interesting way.

```
RBParser >> bracketsOfNode: (RBNode {

left (-> Integer) .

right (->Integer)

}) node
```

Note that in addition to the structural type we have seen, `RBNode` is a nominal type, the type of `node` above limits the value to be not just an instance that understands certain methods, but also requires it to be an instance of `RBNode`.

Another interesting example is the combination of structural type and nominal type `Dyn`:

```
Canvas >> (Self) drawPoint: (Dyn {x (-> Integer) . y (->Integer)}) point

... point x. "safe call"

... point y. "safe call"

... point z. "not an error, considering point to be Dyn"
```

Besides calling `x` and `y`, a call to `z` does not raise an error ahead of execution, because `Dyn` can be casted to allow use of `z`.

2.4.5 Type Aliases

Gradualtalk supports type aliases. This does not expand expressiveness of type itself, but it make reusing existing types convenient.

With type aliases comes the concepts of protocols, which are just type aliases for structural and nominal types.

2.5 $C^\sharp$ 4.0

While we have seen many examples that a dynamically typed language gets gradual type extensions, Bierman, Meijer and Torgersen's work on C$^\sharp$ 4.0 shows us a practical example of statically typed language extends towards gradual typing. This extension makes interoperation between statically and dynamically typed code more convenient and brings in features meant for dynamically typed languages.

2.5.1   Interaction with Dynamic Objects

The following code snippet shows how C$^\sharp$ 3.0 interact with JavaScript:

```
...

ScriptObject map = win.CreateInstance("VEMap", "myMap");

map.Invoke("LoadMap");

...

string latitude, longtitude;

...

var x = win.CreateInstance("VELatLong", latitude, longitude)

var pin = map.Invoke("AddPushpin", x);

pin.Invoke("SetTitle", name);

...
```

In order to interact with JavaScript, method calls are made using string-based interface, which is verbose, fragile to maintain.

With C$^\sharp$ 4.0, the source code can be simplified to:

```
...

dynamic map = win.CreateInstance("VEMap", "myMap");

map.LoadMap();

...

string latitude, longtitude;

...

var x = win.CreateInstance("VELatLong", latitude, longitude)

var pin = map.AddPushpin(x);

pin.SetTitle(name);

...
```

Notice how types are changed to **dynamic** and methods are called as if there are regular objects. These improvements are more than just syntactic: gradual typing is implemented so that typechecking can be involved and this part of code can benefit from it.

2.5.2 Using Features Intended For Dynamic Languages

The extension also makes available some features meant for dynamic languages directly in C$^\sharp$, one example is the use of **ExpandoObject**

```
...

dynamic contact = new ExpandoObject();

contact.Name = "Erik"
```

```
contact.Address = new ExpandoObject();

contact.Address.State = "WA";

...
```

For a statically typed language, to maintain structured object, one either needs to declare a record type with matching structure, or use a dictionary in which one key and one value type should be declared. However, with the introduction of **ExpandoObject**, by simply assigning values to properties of it, they are brought into existence without boilerplate about type casts. This helps in terms of fast prototyping, in which programmers cares more about having a working implementation than maintain type precision.

Chapter 3: The Way Towards Gradual Typing

In previous chapter we have explored the benefits gradual typing can bring in for existing languages. There is no canonical way of supporting gradual typing: it can be extended for either a static typed language or dynamically typed one, and researchers have to make design choices depending on the nature of the language. This chapter takes a step further to discuss the way towards making these extensions. In general, this is not just an effort of putting theories of gradual typing into practical use, but also introductions of new language features that accomplishes the changes and helps development for these languages.

3.1 Extensions to Type Systems

Gradual type system supports both fully static and fully dynamic typing, which requires support for traditional types including basic type, function type and types for structural data and a notion of dynamic type to allow presence of partial type information in the system. Type soundness should be established by a combined work of static and dynamic typechecking. Additionally, some other interesting extensions are also introduced to provide better supports for gradual typing.

This section will separate these extensions into two categories: **Essential extensions** are required for a complete gradual typing support; and **complementary extensions** are those that improves type expressiveness but a type system without them does not compromise completeness - after all, explicit types are optional for gradual typing.

3.1.1 Essential Extensions

We are looking at research works that falls into two different type disciplines: JavaScript, Scheme, Python 3 and Smalltalk are all dynamically typed languages. And $C^\sharp$ 3.0 is statically typed language. While dynamically typed languages share some similarity in their implementations. The one from $C^\sharp$ 3.0 is quite a different story.

3.1.1.1 Extension from Dynamically Typed Languages

The approach used by Safe TypeScript and Reticulated Python is to first perform static typechecking in the presence of dynamic types. If it is successful, programs written in extended languages are then translated into original languages. The translated program could contain runtime checks with some values instrumented, these checks and instruments are either implemented by inserting expressions or statements in place or making calls to external libraries (which is usually called runtime in short) shipped with the extended languages.

This ensures a high compatibility with original languages: source code written

in extended language can be compiled and then used as if it is written in the original language, and thanks to the fact that all existing dynamically typed languages we are discussing have decent supports for modern package managers, including runtime support is as simple as making a dependency declaration to package metadata. Besides, by compiling to the original languages, we can to some extent be confident that the language implementations are independent of specific implementations of the original languages, albeit some language extensions do make assumption about implementations of their original languages.

Gradualtalk is an extension to Pharo Smalltalk. It consists of 3 parts: the core allows representing types in Smalltalk, the typechecker is responsible of performing static typechecks and finally a type dictionary for storing type information.

Typed Scheme takes a similar approach: the implementation is a macro (a source-to-source transformer) that does typechecking and expands source code. This phase either results in static type error being reported, or a working source code in PLT Scheme ready for execution.

3.1.2   Extension from A Static Type System

$C^\sharp$ is the only statically typed language that we have explored in this survey. Similar to Siek's original approach, the compiler assign types to terms using built-in knowledge, explicit type annotations and type inference, then the language is translated into another core language that consists of less constructions and dynamic types are either turned into casts or removed in this process to produce the final

bytecode ready for execution.

Unlike traditional static type systems, $C^\sharp$ 3.0 features bidirectional type system. Such a system have two distinct phases: **type checking** and **type synthesize**. While a type checking phase determines whether a given term can be assigned a particular type, type synthesize phase works out the expected type from context for a given term. These two phases are intertwined to produce a fully typed program.

(TODO: achieve transitivity)

3.1.3 Complementary Extensions

While essential extensions establish the completeness of gradual type systems, complementary extensions are not exactly requirements. After all, programmers can always choose not to use explicitly typed terms and leave types of them inferred or defaulted to dynamic types. This is however unsatisfactory, as using dynamic types often result in a lost of type information in the process. Therefore extra features of type systems are usually introduced to improve type expressiveness and in general make language more convenient to use.

3.1.3.1 Union Types

Practical languages utilize control flows to have different behaviors depending on situations. For a dynamically typed language, an expression can usually return values of various possible types. Typed Scheme and Gradualtalk employ union types to allow preserving more type information.

By notation, an union type $U = \tau_0 \cup \tau_1 \cup \ldots$ typechecks as long as the term in question typechecks with one of τ.

(TOOD: formal)

3.1.3.2 Occurrence Typing

Supporting gradual typing requires having matching types for values and functions in the languages. Most extensions we have discussed have concepts of classes and objects, they are typed either structurally or nominally. However, instead of using classes or objects like other languages do, typical Scheme programs prefers using composed data structures as simple as just pairs, vectors and they are distinguished through testing the shape of the value or symbol comparison. To solve this potential mismatch between values and types, occurrence typing is introduced. The idea is that the difference occurrences of the same variable would have different types depending on the context where it occurs. Types are defined by specifying some boolean-returning unary functions as type discriminators. When such a function `p?` in question returns non-falsy value for a value `v`, the type system will consider `v` to have value `p`. This way, a `if`-expression that depends of value of `(p?  v)`, will have its two branches with different types of `v`, one with type `p` and another with non-`p` type.

3.1.4 Parametric Polymorphism

(TODO: or mark as Typed Scheme-specific?)

3.2   Cast Insertion

Cast insertion provides mechanism for runtime typechecking. During static typechecking, type safety cannot be established in the presence of dynamic types, therefore casts are inserted in corresponding positions of the program. These casts are functions that takes a single value, checks whether the value is of the expected type then either raises type errors or resume the program with casted value.

Unlike static typechecking, cast insertion has large impacts on runtime performance of a program. Many strategies of performing cast insertion are proposed and explored to reduce its overhead.

3.2.1   Object Proxy

Originally described by Siek and Taha, runtime objects are accessed through a proxy to allow field writes and method calls to be monitored.

3.2.2   Shallow Tagging

Shallow tagging approach is used by Safe TypeScript and Reticulated Python as transient mode: instead of tagging objects deeply, they are only tagged up to one layer deep. Accessing deeper into the structure causes invariants to propagate.

3.2.3 Monotonic Approach

Monotonic approach have the effect of locking down objects into precise types. It has the benefit that runtime information can be trimmed to reduce overhead at long run.

3.2.4 Type Dictionary

This approach is used by Gradualtalk, but can also be adapted to other languages of similar nature: in Smalltalk, certain classes are not allowed to be modified because the virtual machine relies on them to implement fundamental behaviors. Therefore language implementor maintains a type dictionary that keeps type information for objects instead of tagging directly on object themselves.

3.3 Dealing with Structural Data

Practical languages come with support for structural data in which multiple values can be combined together to construct one single value, which is usually called an object. These values can be retrieved later using a key or index into that object. Imperative languages would even allow values in an object to be removed or updated at runtime, which complicates some support towards gradual typing. It is such an important concept that all languages we have discussed so far study it to some extent.

3.3.1 Structural Data and Subtyping

3.3.2 Differential Subtyping

The purpose of combining static and dynamic typing is to reduce runtime checks when checks could have been performed once ahead of execution. Differential subtyping by Safe Typing takes this idea and apply it in a finer-grained level.

3.3.3 Support Field Mutation

3.3.4 Object-Oriented Programming

The idea of structural data goes further: object-oriented programming adds methods to objects, and them to be shared by inheritance. This increase flexibility as different runtime behavior are enabled even if method of the same name and type signature is called.

3.3.4.1 Nominal or Structural

Structural type can be useful for languages that uses the practice of duck typing, as only the shape of types are required to be checked. However, nominal types can also find its uses when it comes to object-oriented programming, as structural type does not account for inheritance that object-oriented programming relies heavily on.

3.3.4.2   Self Types

3.3.5   Recursive Data Types

3.4   Eval function

3.5   Other Performance Concerns

3.5.1   Large Arrays

For programs that need runtime typechecks on large arrays, the overhead could be wasteful but significant.

3.5.2   Type Erasure

3.6   Language-specific Challenges

3.6.1   TypeScript

3.6.2   Typed Scheme and Macros

3.6.3   Python

3.6.4   Gradualtalk and Live system

3.6.5 $C^\sharp$ 4.0

Chapter 4:   Related Work

TypeScript implements "occurrence typing" (see "Type Guards and Differentiating Types" of advanced types) and Array as tuple

Chapter 5:   Conclusion